

# Homework 2 Group Part Report

Team: Miguel Fuentes, Dan Cortes, Vineeth Gutta, Praveena Adavikolanu

## Naive Bayes

### Data Import

In [1]:

```
import pandas as pd # this library is used for the dataframe data structure

train1 = pd.read_pickle('data\\dataset 1\\train.pkl')
test1 = pd.read_pickle('data\\dataset 1\\test.pkl')

train2 = pd.read_pickle('data\\dataset 2\\train.pkl')
test2 = pd.read_pickle('data\\dataset 2\\test.pkl')

train3 = pd.read_pickle('data\\dataset 3\\train.pkl')
test3 = pd.read_pickle('data\\dataset 3\\test.pkl')
```

### Model Results

In [3]:

```
import naive_bayes # this holds our implementation of naive bayes

model1 = naive_bayes.NaiveBayes()
model1.train(train1, ['spam', 'ham'])
print(f'Accuracy on data set 1: {naive_bayes.accuracy(model1, test1)}')

model2 = naive_bayes.NaiveBayes()
model2.train(train2, ['spam', 'ham'])
print(f'Accuracy on data set 2: {naive_bayes.accuracy(model2, test2)}')

model3 = naive_bayes.NaiveBayes()
model3.train(train1, ['spam', 'ham'])
print(f'Accuracy on data set 3: {naive_bayes.accuracy(model3, test3)}')
```

Accuracy on data set 1: 0.9602510460251046

Accuracy on data set 2: 0.956140350877193

Accuracy on data set 3: 0.9300184162062615

## Logistic Regression

### Data Import and Formatting

The following block of code is used to process the data into the form which is appropriate for logistic regression

In [4]:

```
from collections import namedtuple

import spam_ham_util # this holds a data formatting utility

DataSet = namedtuple('DataSet', ['X_train', 'X_test', 'Y_train', 'Y_test'])

dataset1 = DataSet(*spam_ham_util.df_to_numeric(train1, test1))
dataset2 = DataSet(*spam_ham_util.df_to_numeric(train2, test2))
dataset3 = DataSet(*spam_ham_util.df_to_numeric(train3, test3))
```

## Hyperparamater Tuning

In the following block we test various values for the regularization constant and see which one results in the best accuracy on a 70/30 split of the training data

In [7]:

```
# This accuracy calculator was not written by us
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split

import log_regression # this holds our implementation of logistic regression

ITERS, LEARNING_RATE = 100000, 0.01
best_reg_val = None
best_acc = 0

for l2_reg in [0, 0.15, 0.3, 0.45, 0.6, 0.75]:
    accs = []
    for dataset in [dataset1, dataset2, dataset3]:
        X_train, X_test, Y_train, Y_test = train_test_split(dataset.X_train, dataset.Y_train, test_size=0.30, random_state=17)

        model = log_regression.log_regression(ITERS, l2_reg, LEARNING_RATE)
        model.train(X_train, Y_train)

        accs.append(accuracy_score(model.predict(X_test), Y_test))
    acc = sum(accs)/len(accs)
    if acc > best_acc:
        best_acc = acc
        best_reg_val = l2_reg
    print(f'Average accuracy with l2_reg = {l2_reg} is: {acc}')
```

```
Average accuracy with l2_reg = 0 is: 0.9226321634204023
Average accuracy with l2_reg = 0.15 is: 0.9351634220573851
Average accuracy with l2_reg = 0.3 is: 0.9397029512345955
Average accuracy with l2_reg = 0.45 is: 0.9392331545391706
Average accuracy with l2_reg = 0.6 is: 0.9321099642025509
Average accuracy with l2_reg = 0.75 is: 0.8805858005107301
```

## Test Set Accuracy

Now that the regularization constant has been chosen we will train on the entire training set and report accuracy on each of the test sets

In [8]:

```
print(f'----- Using l2_reg = {best_reg_val} -----')

for index, dataset in enumerate([dataset1, dataset2, dataset3]):
    model = log_regression.log_regression(ITERS, best_reg_val, LEARNING_RATE)
    model.train(dataset.X_train, dataset.Y_train)

    print(f'Accuracy of dataset{index + 1} is : {accuracy_score(model.predict(dataset.X_test), dataset.Y_test)}')
```

```
----- Using l2_reg = 0.3 -----
Accuracy of dataset1 is : 0.9079497907949791
Accuracy of dataset2 is : 0.9035087719298246
Accuracy of dataset3 is : 0.9539594843462247
```

## Perceptron Algorithm

This code will report the optimal hyperparamaters and the accuracy when trained on all the training data for every dataset individually

In [9]:

```
import perceptron # this is our implementation of the perceptron algorithm

print("DATASET 1-----")
```

```
perceptron.main('data\\dataset 1\\train', 'data\\dataset 1\\test')
print("\nDATASET 2-----")
perceptron.main('data\\dataset 2\\train', 'data\\dataset 2\\test')
print("\nDATASET 3-----")
perceptron.main('data\\dataset 3\\train', 'data\\dataset 3\\test')
```

DATASET 1-----

Data initialized:  
Started training...

Best results from hyperparameter tuning:  
Best Acc: 100.0  
best Epoch: 75  
Best lr: 0.01

RESULTS FROM ENTIRE DATASET:

Learning rate: 0.0100  
Number of epochs: 75  
Emails classified correctly: 443/478  
Accuracy: 92.6778%

DATASET 2-----

Data initialized:  
Started training...

Best results from hyperparameter tuning:  
Best Acc: 100.0  
best Epoch: 75  
Best lr: 0.01

RESULTS FROM ENTIRE DATASET:

Learning rate: 0.0100  
Number of epochs: 75  
Emails classified correctly: 414/456  
Accuracy: 90.7895%

DATASET 3-----

Data initialized:  
Started training...

Best results from hyperparameter tuning:  
Best Acc: 100.0  
best Epoch: 75  
Best lr: 0.01

RESULTS FROM ENTIRE DATASET:

Learning rate: 0.0100  
Number of epochs: 75  
Emails classified correctly: 501/543  
Accuracy: 92.2652%